PROGRESS REPORT · 19 AUGUST 2026

Nexus Agentic OS

Five businesses. One WhatsApp number. An agent that answers in seconds, refuses to invent, and is watched by sixteen checks that never call a model.

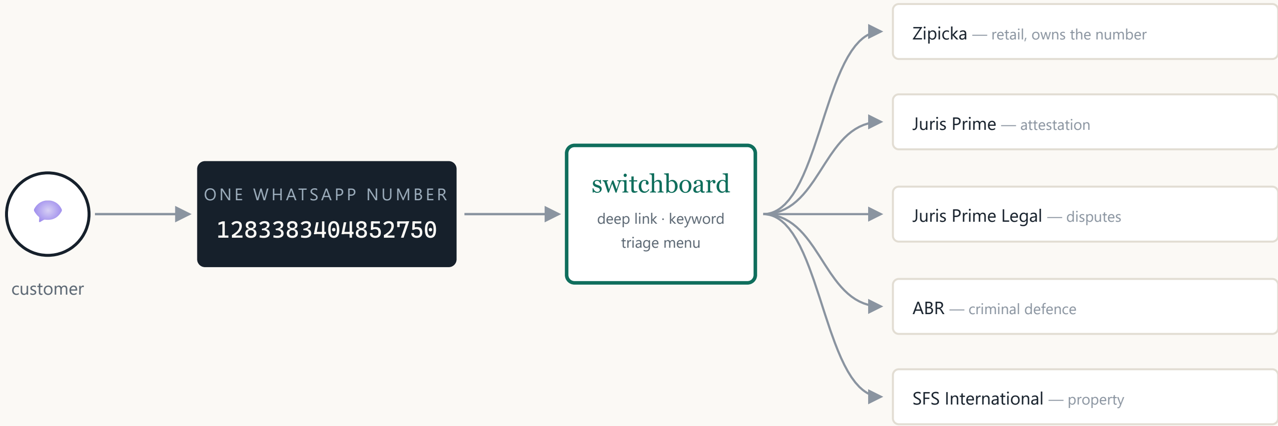

One WABA number, one switchboard, five tenants — and a database that refuses to let one firm read another's customers.

85% PLATFORM COMPLETE	858 TESTS PASSING	10/10 PRODUCTION GATES GREEN	0 STANDING FINDINGS

01 · WHERE IT STANDS

The platform is built and verified. The traffic is not there yet.

Everything below is measured against the live database on 19 August 2026 — not estimated, not remembered. Where a number is small, it is printed small.

What is running

Businesses live	5
Sharing one WhatsApp number	5
Agents configured & active	5
Knowledge sources indexed	65 / 65
Retrievable passages	374
Staff on rota	2

What has actually happened

Conversations, all time	15
Inbound messages	39
Outbound messages	27
Bookings taken	0
Campaigns sent	0
Trading period	1 Aug – 18 Aug

> **Read those two columns together.** The machine is finished and the door is barely open. Thirty-nine inbound messages in eighteen days, across five businesses, is not a product problem — it is a distribution problem, and the single largest lever is four website edits that have been waiting on the owner since 17 August.

Engineering delivered, 20 July – 19 August

284
COMMITS IN 31 DAYS

41,437
LINES OF SOURCE

13,559
LINES OF TESTS

33% of

56
SCHEMA MIGRATIONS

16
MODEL-FREE OPERATORS

10
PRODUCTION GATES

81
TEST FILES / 858 TESTS

34
FINDINGS RAISED, 0 OPEN

02 · PROGRESS BY WORKSTREAM

85% complete — and the derivation, so it isn't just a number

A single percentage is only useful if you can see what went into it. Each workstream is weighted by its share of the product and scored on what is deployed and passing its gate in production — not on what is written.

WORKSTREAM	WEIGHT	COMPLETE
Reply pipeline & switchboard routing, grounding, governance judge, delivery receipts	20%	100
Tenant isolation on a shared number seven defects of one class, found and closed; policies now route-aware	15%	100
Knowledge & retrieval 65 sources, lexical fallback, self-healing re-index after quota failure	12%	100
Monitoring & verification 16 operators, 10 gates — but findings are pull-only: nobody is told	15%	95
Operator console (deck & inbox) every screen live; per-business attribution corrected 19 Aug	13%	90
Revenue features — bookings, campaigns built and never exercised: 0 bookings, 0 campaigns; reply-to-campaign undecided	15%	70
Operations — backup, alerting, handover nightly backup runs to local disk only — offsite copy needs owner credentials	10%	50
Weighted total		85

Weights reflect share of product surface; scores reflect what is deployed and passing its gate in production, not what is written.

What the remaining 15% actually is

REMAINING ITEM	BLOCKED ON	BUILD EFFORT	STATUS
Offsite backups — nightly dump exists, copy is local-disk only	Owner: `rclone` credentials	2 h	WAITING
Four website edits — deep links verified to route	Owner: site access	done	WAITING
ABR office number — a phrase still contains a placeholder	Owner: one number	5 min	WAITING
Alerting — findings are detected and nobody is told	Owner: choose a destination	1 d	DECISION
Campaign replies — a reply arrives with no idea what it answers	Product decision	2 d	DECISION
Aggregate operators — counts still pool under the number's owner	—	0.5 d	QUEUED
Contact ownership — who owns a customer of two firms?	Product decision	1 d	DECISION
First real booking — the diary path has never run for a customer	Traffic	—	WAITING

How long to finish everything

Engineering: about one working week — roughly 4.5 days of build across the seven open items, assuming the two product decisions are answered when asked rather than after.

The critical path is not engineering. Three items need under two hours of the owner's time. Until the four website edits ship, the platform has no way to receive the traffic it was built for — so "finished" arrives in a week, and "working" arrives when the tap is opened. Assumes current pace and no change of scope.

03 · WHAT MAKES IT DIFFERENT

Six things this does that general agentic systems do not

Most "agentic OS" products are a model, a tool loop, and a dashboard. The differences below are structural — they are in the database and the deployment, not in the prompt.

01 Five businesses, one number

Almost every WhatsApp agent is one number per tenant. Here five separate companies — retail, attestation, litigation, criminal defence, property — answer on a *single* WABA number, with a switchboard that routes by deep link, keyword or triage menu. One number's cost and verification overhead carries five brands, and a customer who arrives for one can be handed to another in the same thread.

02 Isolation enforced by the database

Separation is not a `WHERE` clause the code is trusted to remember. Every tenant table carries a row-level security policy, and unscoped queries *throw* rather than silently returning everything. A firm can read the customers it is serving and no one else's — proved on every deploy by a gate that tries to break in.

03 Monitoring with no model in it

Sixteen operators sweep the database every ten minutes in plain SQL — customer waiting, handover abandoned, knowledge broken, delivery failing. They raise *and retract*, so the list reflects now rather than growing forever. Nothing about the watching can hallucinate, drift, or cost a token.

04 It refuses rather than guesses

Answers are grounded in the firm's own indexed material, and a governance judge scores every draft for fabrication *before* it is sent. The forecasting layer declines to project when it has too little history instead of drawing a confident line. In regulated work — legal, attestation — an invented answer is worse than no answer.

05 It won't promise a person who isn't there

Escalation checks the real staff rota and working hours before telling a customer a specialist will follow up, and releases a handover nobody can take rather than leaving the customer muted. This was learned the hard way: four conversations sat silent for sixteen days because the agent promised a colleague who did not exist.

06 Verification that runs against production

Ten gates run against the live system, not a test double — including one that rebuilds the entire schema from scratch in a throwaway database and diffs it against production, and one that proves the running containers were built from the current commit. Each gate exists because something once failed silently.

The defect class this architecture created — and closed

Sharing one number is the platform's leverage and its hardest problem. Because a conversation is owned by the number's holder but served by another firm, a read from the wrong transaction returns *zero rows rather than an error*. Seven separate defects of that one shape were found and fixed. Five had already cost something; the last two were caught before they fired.

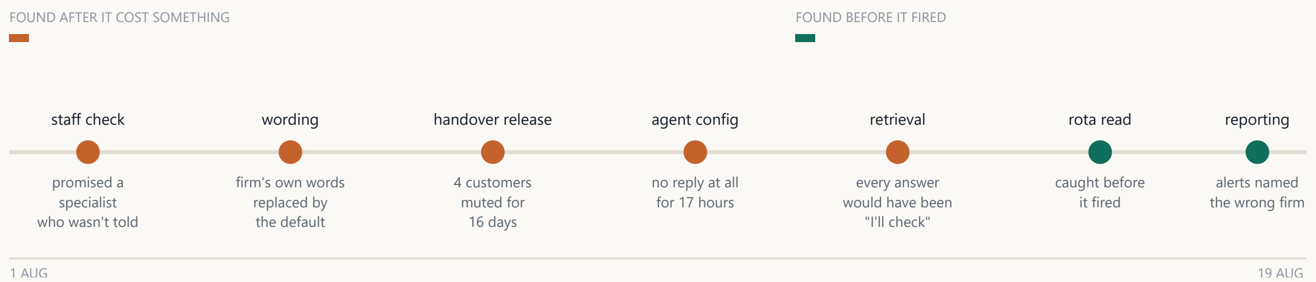

The gate written after instance three is why instances six and seven were caught in advance. The class is now closed at the database, not the call site.

04 · WHAT IT CAN DO FOR THE BUSINESSES

Where the growth actually comes from

These five firms sell time and expertise into a market that messages on WhatsApp at every hour. The platform's value is not "an AI chatbot" — it is that every one of the steps below now happens without a person being awake.

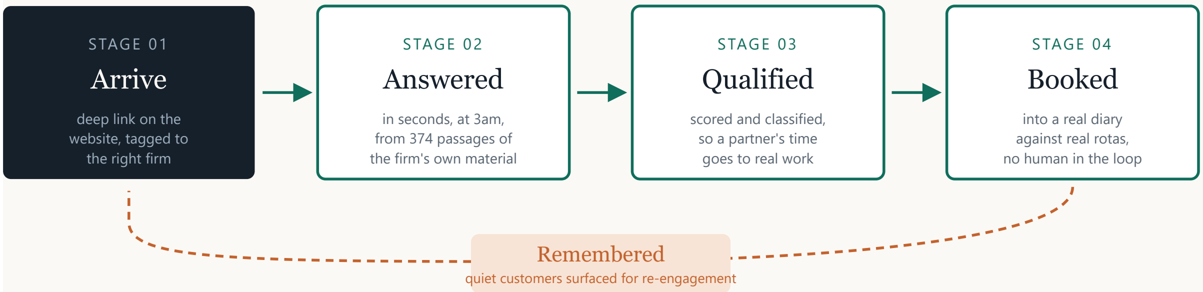

MEASURED TODAY, ALL FIVE FIRMS, 18 DAYS

15	39	27	0
conversations opened	customer messages in	replies out	appointments booked

Stages one to four are built and verified end to end. Stage four has never run for a real customer, because stage one has not been switched on.

The five levers, in order of expected return

LEVER	EFFORT	WHY IT MOVES THE NUMBER
1. Publish the four deep links	owner, 1–2 h	The only thing standing between existing website traffic and a tagged, correctly-routed conversation. Everything downstream is already built.
2. Answer at 3am	live now	Two customers waited 22 and 17 hours for a human. Gulf enquiries arrive across every timezone; first response wins the instruction.
3. Turn talk into diary	live, untested	Booking against real rotas removes the back-and-forth that loses a client between the enquiry and the meeting.
4. Re-engage the quiet	1 decision	The operator already identifies contacts who went quiet. Campaign machinery with delivery receipts is built and has never been sent.
5. Cross-refer inside the group	design	A property client with a lease dispute is a litigation client. One number makes that handoff a routing decision, not a phone call.

What five firms on one number is worth

A single WhatsApp Business number, one verification, one set of templates and one operations console carries five brands. The marginal cost of the sixth business is a row in a table and a knowledge import — not another number, another integration, or another subscription.

The same arithmetic applies to attention. One console shows every firm's waiting customers, broken knowledge and abandoned handovers in one list, ranked by severity and age. One person can hold five businesses.

Stated plainly: none of this has produced revenue yet. Zero bookings and zero campaigns are real numbers, not omissions. The honest claim is that the capability is finished and verified, and the constraint has moved from engineering to distribution.

05 · HONEST LIMITS

What could still go wrong, and what is already guarded

Live risks

RISK	STATE
Backups are local-disk only A nightly dump exists and a floor refuses to run without free space, but nothing leaves the machine. A host failure loses everything.	OPEN
Nobody is told when something breaks Findings are correct and pull-only. One knowledge outage stood 16 hours; a waiting customer stood 22.	OPEN
Free-tier embedding quota Exhausted once on 18 Aug, taking 53 of one firm's 72 passages offline. Now self-heals after 24 h; a paid key removes the ceiling.	GUARDED
Single host One VPS runs database, queue, API, worker and web. No failover.	ACCEPTED
Untested revenue paths Bookings and campaigns pass their gates but have never met a customer.	AWAITING TRAFFIC

What is already guarded

858 tests
every commit — a hook blocks the commit if one is red

10 gates
every deploy — run against production, not a test double

16 operators
every 10 minutes — plain SQL, raise and retract

6 job heartbeats
an outside probe — catches the watcher going quiet

Each layer watches something the layer above it cannot see. The bottom one exists because the sweep that finds everything else cannot report its own silence.

In one paragraph

Nexus Agentic OS is finished to 85% and verified against production on every deploy. It does something most agentic systems do not attempt — five separate companies trading through one WhatsApp number, isolated by the database rather than by convention, watched by monitoring that cannot hallucinate. About a week of engineering remains. The larger gap is that thirty-nine customer messages have reached a system built to handle thousands, and closing it starts with four links on four websites.

Nexus Agentic OS — progress report, 19 August 2026. Every figure in this document was read from the live production database or the repository on the day of writing. Counts of conversations, messages, bookings and campaigns are all-time totals since 1 August 2026. The 85% figure is derived from the weighted table on page 3 and is an estimate; the measurements it rests on are not.